Google Cloud H2S

AI Builder Cup

Build in JAPAC.
Build for the world.



Team Details

a. Team name: **Koffee after KitKat**

b. Team leader name: **Siddharth Bhardwaj**

Product: **aftercall**

You keep your conversations, we handle the rest.

c. Problem Statement

Retail & Commerce – consumer intelligence

Retail and e-commerce brands talk to customers all day, in support calls, sales calls, WhatsApp and chat, but listen to almost none of it. Teams read a small sample by hand, so why customers buy, return, complain or churn, which products they ask for, and how agents handle them stays buried in transcripts and recordings. aftercall turns every conversation into structured consumer intelligence that teams act on once the call is done.

Brief about the idea

aftercall asks every conversation the same questions

- You define trackers: groups of signals, each a Choice, a Yes/No or a Score.
- Upload transcripts (CSV, JSON, JSONL, TXT, VTT, SRT, any layout) or call recordings (mp3, wav, m4a, ogg, webm, flac, up to 30 MB).
- TypeSafe Jev answers every signal on every call with a confidence score; low-confidence answers go to human review.
- Ask alteria (Gemini) answers questions across all analysed calls, with cited calls.

Who it's for

1. D2C and e-commerce brands
2. Marketplaces
3. Retail chains
4. Quick-commerce support teams
5. AI voice agents for retail

100%

of conversations read, not a sample

Any tool

not tied to a call or chat product

Hindi · Hinglish · English

transcribed in the original language

Opportunities

1 How is it different?

- Works on transcripts and recordings from any tool; no call platform to adopt
- You write the questions; there is no fixed taxonomy
- A confidence score on every answer, not a black-box summary

2 How does it solve the problem?

- Reads 100% of conversations instead of a hand-read sample
- Structured answers show trends by channel, agent or campaign
- Low-confidence answers go to people; Ask alteria explains the why with cited calls

3 USP

- The same questions on every call, at a fraction of an LLM's cost
- Jev routes each request to the cheapest Gemini tier that is enough
- Cost is shown before anything runs; over-limit calls are flagged, not run

You keep your conversations, we handle the rest.

List of features offered by the solution

1 Trackers and signals

Choice, Yes/No or Score questions, grouped by topic

2 Any transcript or recording

Any layout; Gemini transcribes audio with speakers labelled

3 Jev answers + confidence

Every call, every signal; low confidence goes to review

4 Conditional trackers

Run a tracker only if another answer matches

5 Insights across uploads

Trends, breakdowns by channel, agent or metadata, movers

6 Ask alteria

Gemini answers across all calls with cited calls; drafts trackers

7 Dashboards and CSV export

Charts you check every week, shareable and exportable

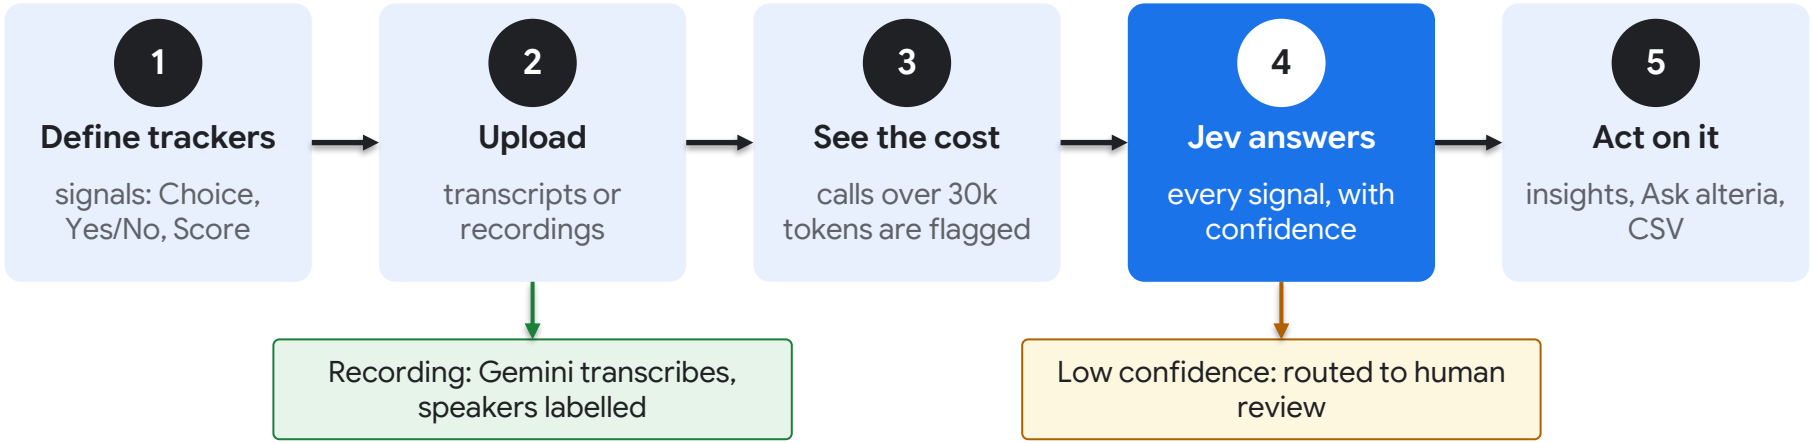
8 People, roles, projects

Workspaces with owners, members and viewers

9 Cost shown first

Estimate before any run; calls over 30k tokens flagged

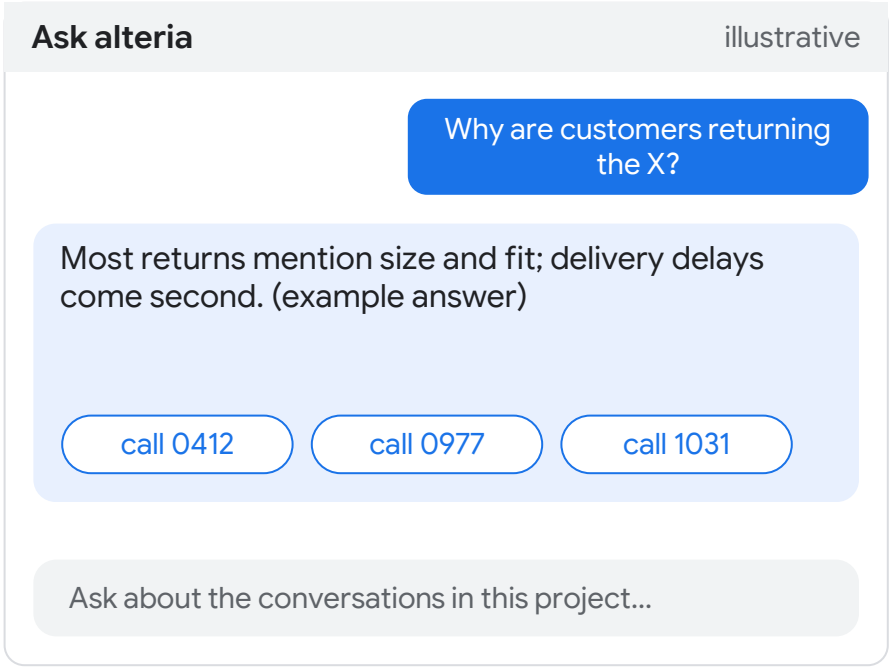
Process flow diagram



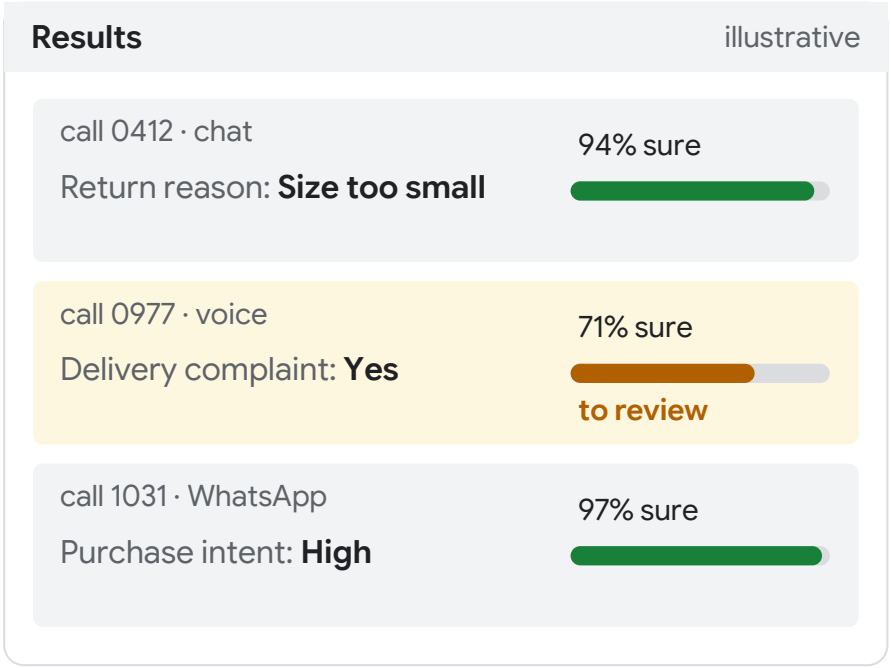
Illustrative retail signals a tracker could ask on every conversation

return / refund reason	purchase intent	product interest	competitor mentioned
price objection	delivery complaint	churn risk	sentiment / frustration
right resolution offered	upsell attempted	policy misstated	asked for a human

Wireframes / mock diagrams of the proposed solution

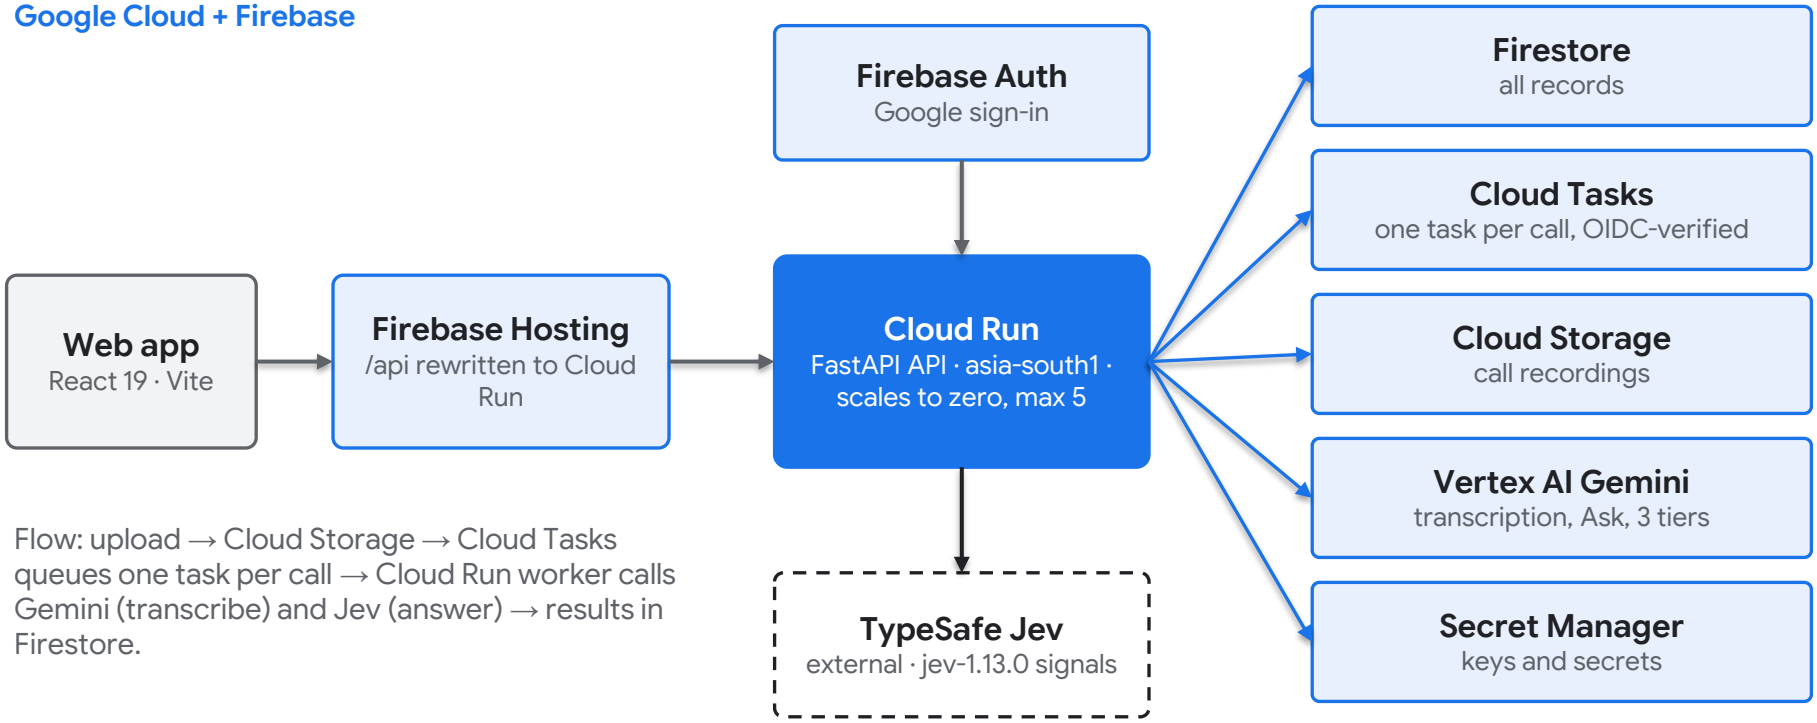


Ask in plain English; answers cite the calls behind them



Every answer shows its confidence; uncertain ones are flagged

Architecture diagram of the proposed solution



Flow: upload → Cloud Storage → Cloud Tasks queues one task per call → Cloud Run worker calls Gemini (transcribe) and Jev (answer) → results in Firestore.

Technologies used in the solution

Google Cloud

- Gemini on Vertex AI
- Cloud Run
- Firestore
- Cloud Tasks
- Cloud Storage
- Firebase Authentication
- Firebase Hosting
- Secret Manager

Models

TypeSafe Jev (jev-1.13.0): signals. Gemini 3.8 Flash: standard + transcription; 3.5 Flash-Lite: simple; 3.1 Pro preview: complex

Web app

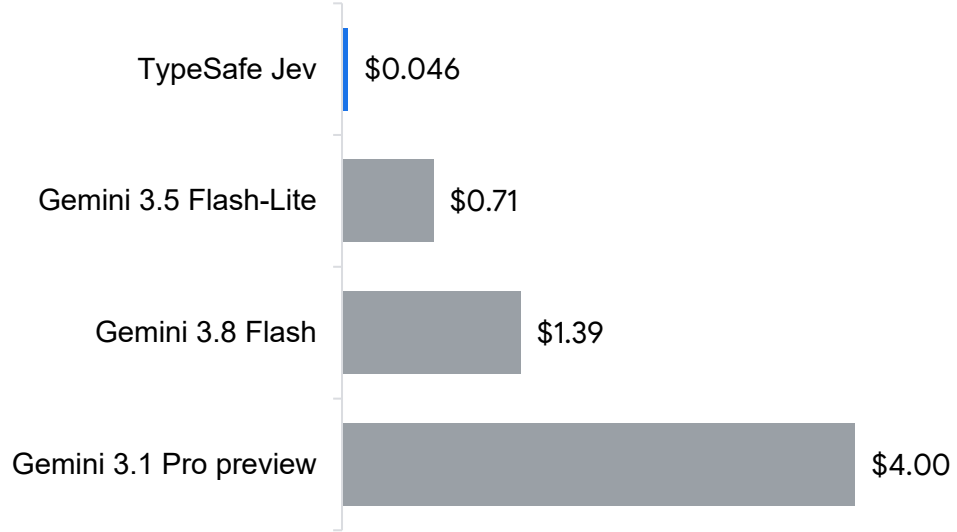
React 19 · Vite · Tailwind · TypeScript

Backend

Python 3.12 · FastAPI · pydantic

Estimated implementation cost

Analysing 1,000 calls, 5 signals each (est.)



Same ~1,100-token call with ~150 output tokens on Gemini. List prices; estimates.

≈ $0.00005

per 5-signal call on Jev (~1,100 input tokens)

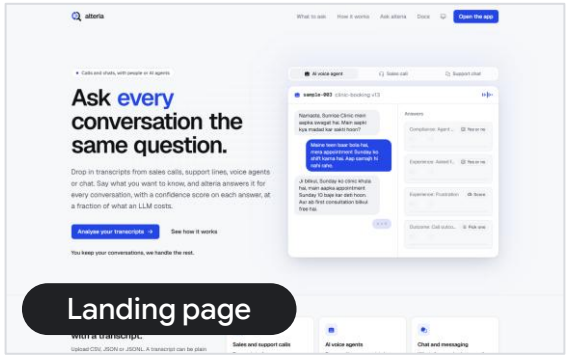
≈ $0.02

to transcribe a 10-minute recording on Gemini 3.8 Flash

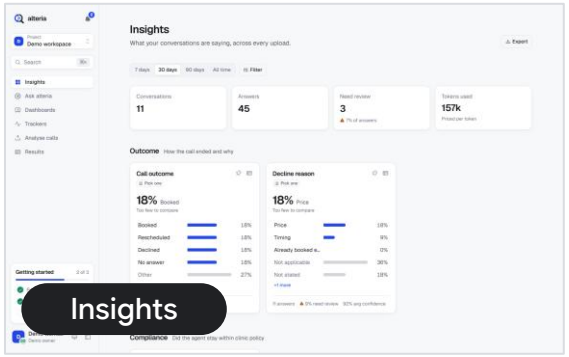
≈ $0

Google Cloud infra at demo scale (free tiers)

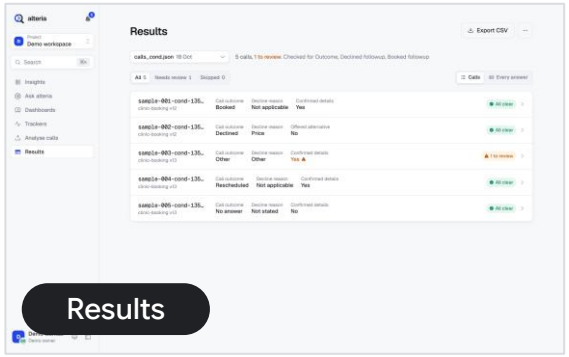
Snapshots of the prototype



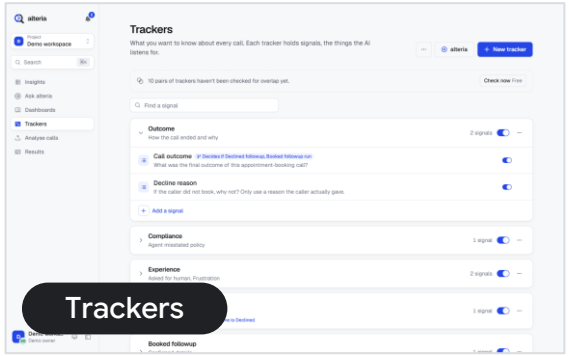
Landing page



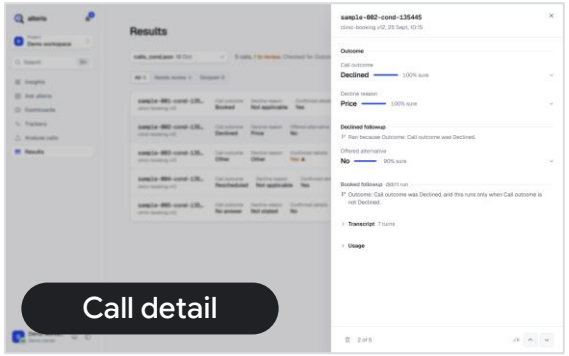
Insights



Results



Trackers



Call detail

Real screens from the running build

Shown on a sample dataset (clinic appointment-booking calls). The same trackers carry over to retail signals such as returns, delivery and churn.

Prototype performance report / benchmarking

25/25

answers matched hand labels (100%): 5 sample calls × 5 signals, live Jev

~6 s

to analyse 6 calls end to end

254

passing backend tests

5/5

conditional trackers routed correctly

~1,100

billed input tokens for a typical call, 5 signals

29.6k → 5

a ~29.6k-token call auto-split into 5 Jev requests; a 32k call flagged, not sent

Efficiency: Insights read only call documents with a 30 s shared cache, about 20× fewer Firestore reads per view (≈600k → ≈10–30k at 10k calls × 10 signals). Low Gemini thinking on simple tiers and transcription.

Confidence: sample answers mostly 0.84–1.00; low-confidence ones routed to review.

Additional details / future development

1

Long calls

Beyond the 30k-token limit

2

Live analysis

Streaming, mid-call signals

3

Connectors

Twilio, Exotel, Genesys, WhatsApp, CRM

4

Regional languages

Hindi and regional-language benchmarks

5

Cheaper transcription

Flash-Lite after quality tests

6

Billing

Per-token prepaid credits, Pro plan back on

7

Alerts

On red-flag signals, e.g. churn risk, policy misstated

8

Multi-region

Run closer to customers

Links to our GitHub repository, demo video and final product

1. GitHub public repository

   [GitHub link]

2. Demo video (3 minutes)

   [Demo video link]

3. Final product

   https://gen-lang-client-0902222994.web.app

You keep your conversations, we handle the rest.

You keep your conversations, we handle the rest.

1. **Every conversation**
   asked the same questions
2. **Every answer**
   with a confidence score
3. **Every insight**
   from trends to Ask alteria

aftercall · Koffee after KitKat · gen-lang-client-0902222994.web.app

Google Cloud H2S

AI Builder Cup

Build in JAPAC.
Build for the world.

Thank you

